Arithmetic, Answer Formatting, and Generalization in a Tiny Language Model Trained with GRPO

Mohamed Sobhy

Living experiment report — October 2026

Abstract

This project investigates direct reinforcement learning on a pretrained 0.8-billion-parameter language model for Countdown arithmetic. A pilot run improved strict greedy accuracy from 0/256 to 85/256 (33.2%) on a fixed held-out set after 300 optimizer steps. This report distinguishes measured findings from planned experiments. A post-hoc audit using a conservative declared-expression parser recovers the same correct counts when the answer-tag requirement is relaxed. This supports investigating arithmetic task improvement, while not establishing general reasoning or fully identifying the role of formatting. Output-budget, multi-seed, reward-ablation, and distribution-shift experiments are tracked below as they are completed.

1 Question and experimental scope

When does direct GRPO improve arithmetic task performance, and how much of an apparent improvement is answer-format compliance? The base checkpoint is `Qwen/Qwen3.5-0.8B-Base`; the dataset is `Jiayi-Pan/Countdown-Tasks-3to4`. The experiment omits supervised fine-tuning. It does not train a model from scratch and does not claim discovery of general reasoning. Model and dataset revisions, code, and outputs are retained in the public repository.

2 Pilot protocol

The text backbone is trained with BF16 rank-16 LoRA on one RTX 4090 Laptop GPU with 16 GB VRAM. GRPO uses eight candidates per puzzle, 32 completions per optimizer step, a learning rate of 10^{-5}, no KL penalty, and a 256-token completion limit. The 300-step run used seed 42 and took 5,135.9 seconds of training time (85.6 minutes), excluding setup and evaluation.

Puzzles are deduplicated by $(\text{target}, \text{sort}(\text{numbers}))$. A seeded hash partition reserves 2% for evaluation before subset selection. The available training subset contains 20,000 puzzles; evaluation uses 256 fixed puzzles. Both base and adapter evaluations use greedy decoding with the same output limit. No evaluation-based checkpoint selection was performed.

The strict reward is binary: one final `<answer>` block must contain a valid arithmetic expression using each supplied number exactly once. The verifier supports binary $+, -, \times, /$, parentheses, and exact rational arithmetic. It rejects number reuse/omission and unsafe syntax. Negative and fractional intermediate values are allowed. No formatting bonus was used.

3 Pilot results

Metric on 256 held-out puzzles	Base	GRPO adapter
Strict correct	0 (0.0%)	85 (33.2%)
Strict valid expressions	11 (4.3%)	90 (35.2%)
Declared-expression diagnostic correct	0 (0.0%)	85 (33.2%)

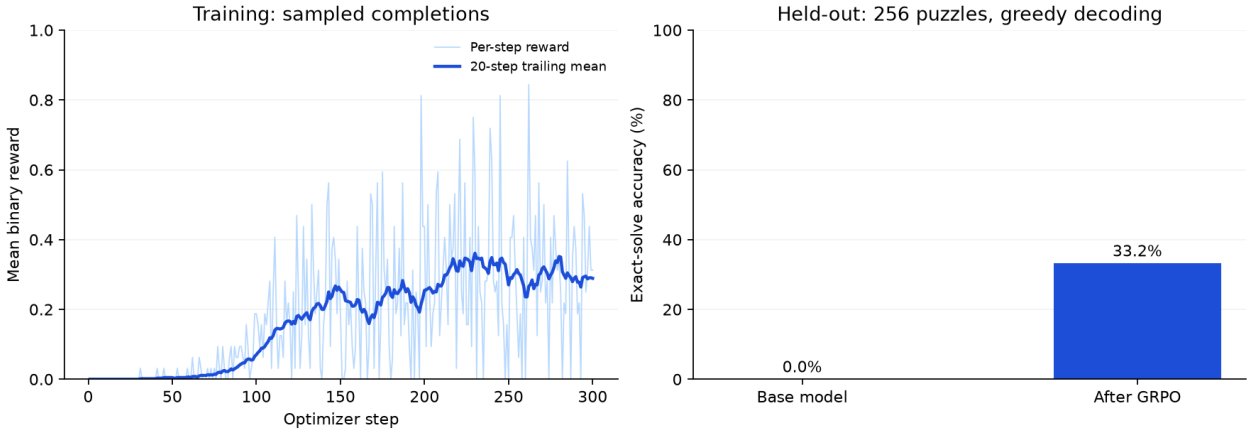


Figure 1: Pilot stochastic training reward and separate greedy held-out accuracy. These use different puzzles and decoding protocols; validation was not measured at every training step.

The first positive training reward occurred at step 31. Peak PyTorch-allocated memory was 2.72 GiB, which excludes reserved memory and other GPU allocations. In the final 30 training steps, 59.9% of sampled completions reached the output cap. Historical evaluation reports do not record termination token IDs, so their truncation status is unknown rather than inferred from decoded text.

4 Format-independent diagnostic: declared-expression-v1

This diagnostic is post-hoc and is not used as a training reward. Extraction is deterministic and independent of the target:

1. If answer tags occur, require exactly one complete block and extract it. A wrong tagged answer never falls back to an earlier expression.
2. Otherwise, select the last standalone arithmetic line, allowing an optional "Answer:" or "Final answer:" prefix and simple math/code delimiters.
3. For an equation, evaluate only the leftmost expression. Do not trust its claimed right-hand value or search intermediate calculations for a success.
4. Apply bounded AST interpretation, exact rational evaluation, and multiset validation of input numbers. Reject unsupported operations and syntax.

This conservative parser can miss correct answers expressed in prose or other notation. Therefore equal strict and diagnostic counts do not prove formatting had no causal role in training. They show that this specified relaxation did not recover additional correct answers from the archived generations.

Diagnostic outcome	Base	Adapter
Correct	0	85
Wrong numerical value with valid numbers	90	5
Number multiset mismatch	152	121
Invalid syntax	10	19
No declared expression	4	24
Ambiguous or incomplete answer tags	0	2

These are mutually exclusive diagnostic categories. Number omission/reuse details are retained per answer; truncation is a separate flag when token metadata exists. All 512 historical outputs were re-scored, and strict scores matched their original records. The same unique holdout puzzles were verified in both reports, with no intersection with the selected training subset. An independent exact subset-dynamic-programming solver found a valid witness for all 256 puzzles under the all-inputs-once rules. Unsolvable holdout items therefore do not explain the errors. Witnesses are archived separately and never included in model prompts. An agent review of the first two records per observed outcome category (20 records total) found agreement with the specified extraction rule; this is not a blinded human annotation study. Prose-heavy and incomplete outputs illustrate the parser's coverage limitations.

5 Controlled study design and current status

The study keeps the original strict metric alongside the diagnostic. The budget grid is complete; remaining comparisons are not presented as completed results:

- Matched base/adapter output budgets of 256, 512, and 1024 tokens; record actual generated lengths, EOS status, truncation, and generation time.
- Binary correctness versus a 0.05 bonus for valid incorrect expressions, with training seeds 42, 43, and 44 and a fixed split seed of 42.
- Three-number-only training with disjoint three- and four-number evaluation.
- Separately labeled, deterministic, solver-verified suites with larger inputs and out-of-range targets; solutions are never included in prompts.
- A comparison demo with exact verifier feedback and archived examples.

Per-seed results and held-out uncertainty will be reported separately. Four distribution-shift/arity suites have been frozen before querying trained models.

Uncertainty and paired analysis. Individual checkpoint accuracies use 95% Wilson score intervals, following the NIST formulation. For the pilot's 85/256 score this interval is 27.7–39.2%. It describes puzzle-sampling uncertainty under a binomial model, not training-seed variability or uncertainty about arbitrary out-of-distribution tasks. Comparisons match canonical puzzles and report gained, lost, jointly correct, and jointly incorrect answers. Exact two-sided McNemar tests are exploratory, with unadjusted p-values; we do not use the grid to claim a confirmatory significance result. The controlled-study aggregator requires all six runs to finish and checks their full training configuration, all 300 optimizer-step records, 20,000 unique training puzzles, identical ordered train/holdout lists, absence of overlap, adapter paths, and recomputed evaluation scores. It reports per-seed accuracies and sample standard deviations, together with paired reward differences for each seed. Repeated evaluation of the same holdout is not pooled into 768 independent puzzles. Synthetic manifest tests reject unfinished runs and altered scores. Hardware/package equality must additionally be checked in the linked W&B provenance before final claims.

Completed output-budget experiment. The base model solved 0/256 at all three caps. The adapter solved 85/256 (33.2%) at 256 tokens and 86/256 (33.6%) at both 512 and 1024 tokens. Both longer caps gain the same single solution and lose none of the original successes. Strict and diagnostic scores coincide in every cell. The adapter's evaluation time increased from 184.6 to 358.6 to 710.0 seconds, while cap-reaching outputs declined only from 145 to 140 to 139 out of 256. Thus quadrupling the output cap costs approximately 3.85 times the evaluation time for one additional solution on this checkpoint and suite. These timings include batched generation and scoring, excluding model loading; they are not single-request latency. We do not infer that longer outputs never help other models or tasks. The analysis run archives every completion and source

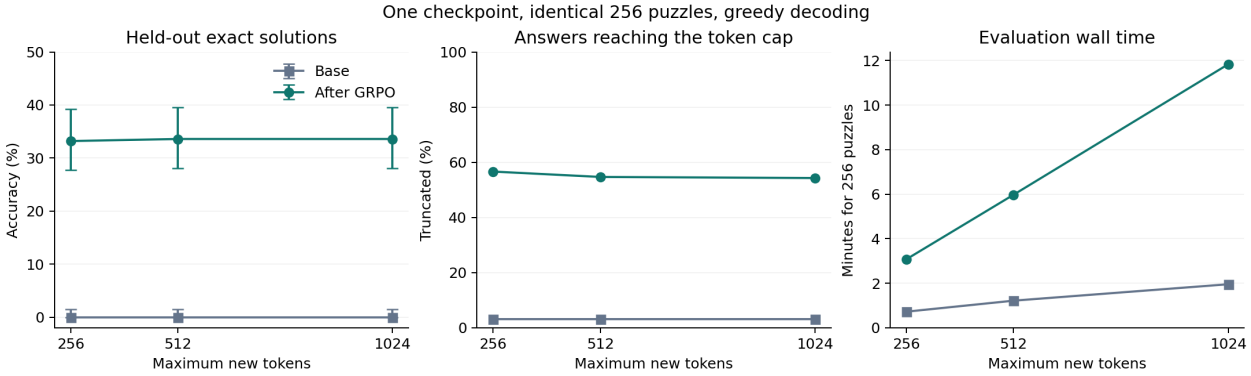


Figure 2: Complete matched budget grid. Error bars are 95% Wilson intervals over puzzles for one checkpoint, not training-seed variation.

hash. The earliest two reports predate the optional `split_seed` and `train_num_count` fields; absent values were normalized to their unchanged `None` defaults. All other settings matched, canonical puzzle sets were identical, and scores were recomputed before aggregation.

6 Pilot generalization on frozen suites

The original mixed-arity pilot was evaluated with 256-token greedy decoding on four solver-verified, frozen suites, each containing 256 puzzles and no overlap with its selected training puzzles. The base model scored zero strict solutions on every suite. The trained checkpoint's results were:

Suite	Solved	Accuracy (95% Wilson)	Truncated
Three-number holdout	146/256	57.0% (50.9–62.9)	91/256
Four-number holdout	41/256	16.0% (12.0–21.0)	184/256
Inputs 101–200	61/256	23.8% (19.0–29.4)	165/256
Targets 101–200	123/256	48.0% (42.0–54.2)	122/256

Diagnostic scores equal strict scores for the trained checkpoint. For the base, diagnostic extraction recovers one solution in each shifted suite, versus zero strict solutions; it recovers none in the two arity holdouts.

The 41-percentage-point three/four-number gap indicates strong sensitivity to task complexity under this evaluation. Number-multiset errors account for 176/256 four-number outputs. This describes observed errors, not a causal explanation of the gap. The arity holdouts may overlap the original pilot evaluation, so they are not independent replications and should not be pooled with

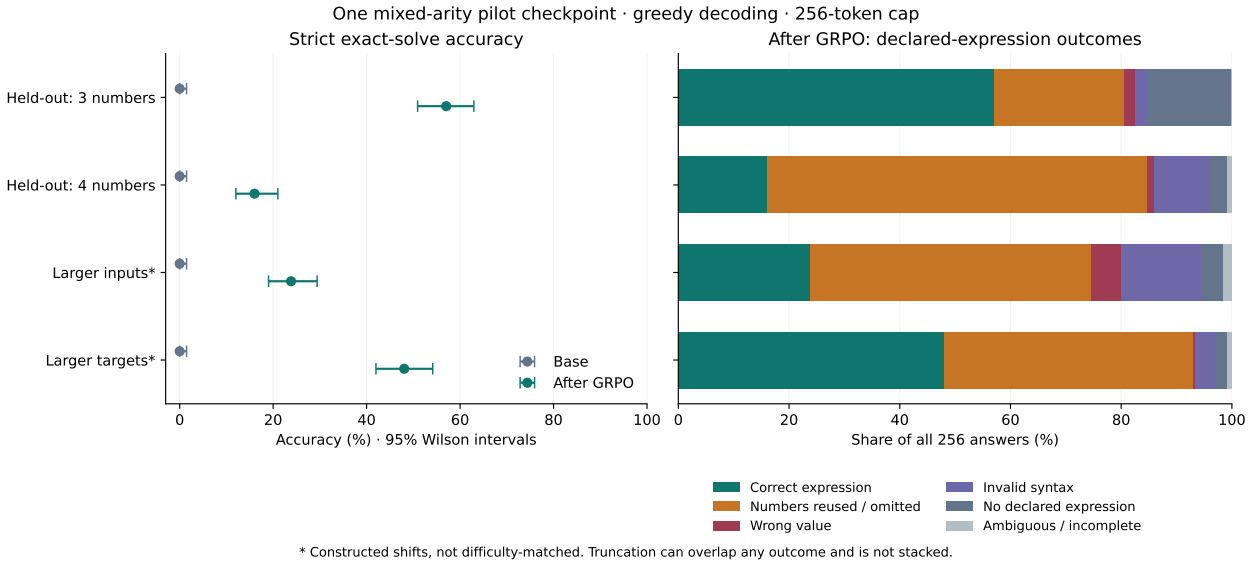


Figure 3: Strict solve accuracy and diagnostic outcomes for the single mixed-arity pilot. Error bars describe per-checkpoint puzzle uncertainty, not seed variation. Truncation overlaps diagnostic categories and is reported separately in the table.

that test set. The shifted suites balance three/four-number inputs and sample targets uniformly from reachable integers within their designated ranges. They are constructed distributions, not difficulty-matched versions of the original dataset. In particular, the 48% larger-target score does not establish superior out-of-distribution reasoning. Full per-answer evidence, hash validation, and paired model comparisons are in the generalization analysis. Training only on three-number puzzles and testing transfer remains a separate, ongoing experiment. Its prepared aggregator requires all three training seeds and all 15 evaluations, checks that every training puzzle contains exactly three numbers, verifies all 300 steps and identical ordered splits, and re-scores every answer against the frozen suites. An end-to-end synthetic test rejects four-number training contamination before publishing output. Per-seed comparisons with the single mixed-arity pilot will be descriptive; they are not a balanced multi-seed ablation of the training distribution.

7 Kernel validation and controlled-study backend

The kernel investigation archives the probes, successful five-step trial, and interrupted fallback evidence. After the hardware trial, FLA 0.5.2 was installed in an isolated directory and compared with unwrapped Transformers reference kernels at representative model shapes. A probe-only positional-argument error initially passed `beta` as `gk` to the recurrent function; inspection of the actual signatures and Transformers calls identified and corrected it. This was not a training failure. The corrected probe found about 3.27× chunk-forward/backward and 2.11× recurrent-forward microkernel speedups. Chunk output relative L2 difference was 0.44%, and gradient differences were 0.27–0.53%; the single recurrent output matched exactly. Returned states were checked for finiteness, not long-horizon equivalence. Compilation is excluded and concurrent GPU work limits the precision of these microtimings.

A separate five-step full-training trial completed in 203.3 seconds, with steps 2–5 averaging 37.63 seconds versus the fallback's approximately 50.8 seconds. All five sparse rewards were zero,

so the test establishes operational viability, not learning quality. The original training process resumed automatically after a recorded 266.62-second suspension. On this evidence, the fallback attempt was subsequently stopped after 41 steps and archived, and all six controlled reward runs were restarted consistently with the pinned FLA backend in `controlled-l4-fla-v1`. Neither short trial enters the completed reward comparison. Original pilot and local transfer runs retain their fallback backend; this distinction must be retained in interpretation and reproduction.

8 Interactive comparison demo

A local browser interface exposes all 256 archived evaluation pairs, including failures, with filters for outcomes and number-multiset errors. An exact answer checker accepts user expressions and shows strict success, diagnostic success, the extracted expression, exact value, and missing or reused inputs. Archived mode uses only Python's standard library and requires no model download or GPU. An optional live mode loads the pilot LoRA adapter and compares greedy generation with the adapter enabled and disabled; its GPU smoke test remains pending. The server binds to loopback and serializes live inference requests.

Chromium checks verified all 256 examples, 85-success/171-failure filtering, interactive answer checking, disabled generation in archived mode, no JavaScript errors, and mobile layout without horizontal overflow at 390px width. HTTP tests also check invalid inputs and cross-origin request rejection. Browser checks are reproducible with `scripts/check_demo_browser.py`; these do not substitute for testing live model inference. Single-puzzle generation can differ from the historical batched evaluation, and custom puzzles are not guaranteed solvable.

9 Provenance and experiment ledger

The pilot used clean source commit `0a432a272591f7683b711a9c867c01e7d939fca1`. W&B contains source snapshots, configs, package versions, puzzle splits, checkpoints, and generation tables:

- Pilot baseline
- Pilot training
- Pilot final evaluation
- Final pilot adapter artifact
- Study workspace
- Archived-output diagnostic audit

2026-10-10. Inspected the current repository, pilot artifacts, local GPU, and Overleaf access. Recorded the approved research plan. Implemented and tested the declared-expression diagnostic, re-scored archived outputs, and added token termination metadata and output-budget overrides to evaluation. Twenty-five CPU tests initially passed; solver coverage increased this to 29 tests. Created and compiled this Overleaf report, verified all holdout puzzles are solvable, and reviewed a deterministic diagnostic sample. The matched output-budget grid is complete, as is the pilot generalization grid. The initial L4 fallback attempt averaged approximately 50.5 seconds per step over its first five steps; the subsequent kernel validation and consistent FLA restart are described above. These are preliminary throughput measurements, not completed study runtimes. Statistical analysis, demo checks, and schema compatibility validation bring current CPU coverage to 48 passing tests, including archive-integrity checks. An automatic collector preserves each completed remote job and exports W&B hardware and software provenance. Analysis requires matching clean source, Python and package environments, and the pinned FLA backend on a single L4.

The dedicated VM is released only after all six completed jobs are downloaded and validated; aggregate analysis then runs locally. These checks do not imply that the pending experiments have completed. An A100 hardware trial completed three steps at 49.33 seconds per step, compared with the L4's then-current 50.79 seconds per step. This approximately 3% gain did not justify switching the controlled grid. The A100 run was intentionally interrupted, its complete partial evidence archived, and its VM released. It is excluded from completed training comparisons. These few steps are a throughput trial, not a reliable general hardware benchmark.

10 Limitations

One seed and 256 puzzles do not establish robust generalization. Training-dataset deduplication cannot rule out exposure in pretraining. Strict success conflates arithmetic correctness with protocol compliance; the diagnostic is an imperfect post-hoc measure, not a causal decomposition. Long answers may contain useful work but still fail due to truncation. The completed output-budget finding applies to this checkpoint and test set; cross-seed and OOD findings remain pending.