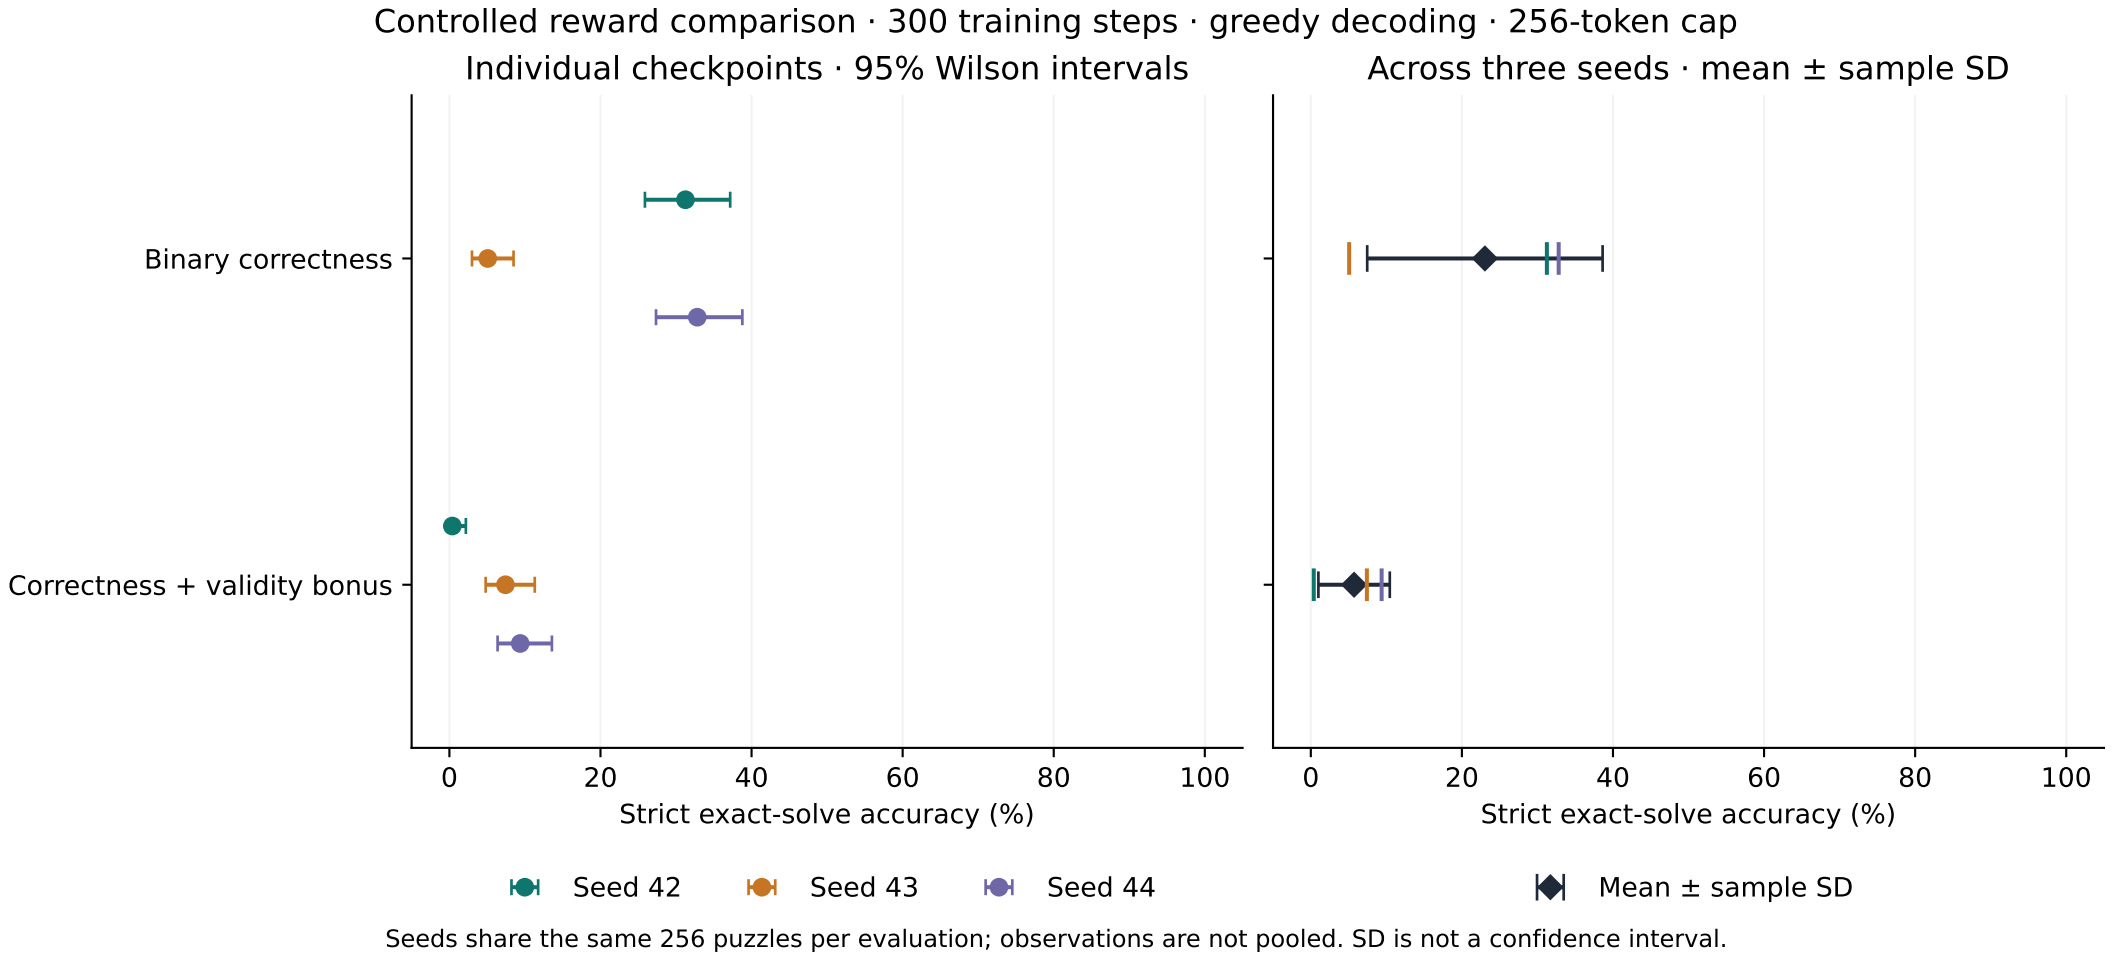
Controlled reward comparison · 300 training steps · greedy decoding · 256-token cap
Individual checkpoints · 95% Wilson intervals
Across three seeds · mean ± sample SD
Binary correctness
Correctness + validity bonus
0
20
40
60
80
100
Strict exact-solve accuracy (%)
Seed 42
Seed 43
Seed 44
Mean ± sample SD
Seeds share the same 256 puzzles per evaluation; observations are not pooled. SD is not a confidence interval.